



# **BIG DATA ANALYTICS/MANAGEMENT CS 6350**

---

Pig Latin



# What is Pig?

---

- Pig is a platform for analyzing large data sets.
- Pig's language, Pig Latin, lets you specify a sequence of data transformations such as merging data sets, filtering them, and applying functions to records or groups of records.
- At Yahoo! 40% of all Hadoop jobs are run with Pig.



# Various Modes

---

- You can run Pig (execute Pig Latin statements and Pig commands) using various modes.

	<b>Local Mode</b>	<b>Mapreduce Mode</b>
<b>Interactive Mode</b>	yes	yes
<b>Batch Mode</b>	yes	yes



# Execution Modes using the **pig** command

---

- **Local Mode**

- Local host and file system.
- /\* local mode \*/
- \$ pig -x local ...

- **Mapreduce Mode**

- Access to a Hadoop cluster and HDFS installation.
- /\* mapreduce mode \*/
- \$ pig ...  
or
- \$ pig -x mapreduce ...



# Interactive Mode **Local** **Mode**

---

- using the Grunt shell.
- Can be invoked using:
  - `$ pig -x local`  
... - Connecting to ...
  - `grunt>`
  - `grunt> A = load '/etc/passwd' using PigStorage(':');`
  - `grunt> B = foreach A generate $0 as id;`
  - `grunt> dump B;`



# Interactive Mode

## Mapreduce Mode

---

- using the Grunt shell.
- Can be invoked using:
  - `$ pig -x mapreduce`  
... - Connecting to ...
  - `grunt>`  
or
  - `$ pig`  
... - Connecting to ...
  - `grunt>`



# Batch Mode

---

- `/* id.pig */`
- `A = load '/etc/passwd' using PigStorage(':'); -- load the passwd file`
- `B = foreach A generate $0 as id; -- extract the user IDs`
- `store B into 'id.out'; -- write the results to a file name id.out`



# Batch Mode

---

- **Local Mode**

- \$ pig -x local id.pig

- **Mapreduce Mode**

- \$ pig id.pig

- or

- \$ pig -x mapreduce id.pig





# Pig Scripts

---

- To place Pig Latin statements and Pig commands in a single file.
- Using the \*.pig extension is good (please do it for HW too).



# Pig Latin Statements

---

- `OutputRelation = InputRelation`
- A relation is a bag.
- A bag is a collection of tuples.
- A tuple is an ordered set of fields.
- A field is a piece of data.



# Debugging

---

- Use the DUMP operator to display results to your terminal screen.
- Use the DESCRIBE operator to review the schema of a relation.
- Use the EXPLAIN operator to view the logical, physical, or map reduce execution plans to compute a relation.
- Use the ILLUSTRATE operator to view the step-by-step execution of a series of statements.



# Debug: dump

---

- A = LOAD 'student' AS (name:chararray, age:int, gpa:float);
- DUMP A;  
(John,18,4.0F)  
(Mary,19,3.7F)  
(Bill,20,3.9F)  
(Joe,22,3.8F)  
(Jill,20,4.0F)



# Debug: describe

---

- `grunt> A = load '/home/kma041000/pig/input' as (line:chararray);`
- `grunt> describe A;`  
`A: {line: chararray}`
- `grunt>`



# Debug: explain

grunt> explain A

```
#-----  
# New Logical Plan:  
#-----  
A: (Name: LOStore Schema: line#3:chararray)  
|  
|---A: (Name: LOForEach Schema: line#3:chararray)  
|   |  
|   | (Name: LOGenerate[false] Schema: line#3:chararray)ColumnPrune:InputUids=[3]ColumnPrune:OutputUids=[3]  
|   | |  
|   | | (Name: Cast Type: chararray Uid: 3)  
|   | | |  
|   | | |---line:(Name: Project Type: bytearray Uid: 3 Input: 0 Column: (*))  
|   | | |  
|   | | |---(Name: LOInnerLoad[0] Schema: line#3:bytearray)  
|   | | |  
|   | | |---A: (Name: LOLoad Schema: line#3:bytearray)RequiredFields:null  
|   | |  
|   |  
|  
#-----  
# Physical Plan:  
#-----  
A: Store(fakefile:org.apache.pig.builtin.PigStorage) - scope-5  
|  
|---A: New For Each(false)[bag] - scope-4  
|   |  
|   | Cast[chararray] - scope-2  
|   | |  
|   | |---Project[bytearray][0] - scope-1  
|   | |  
|   | |---A: Load(/home/kma041000/pig/input:org.apache.pig.builtin.PigStorage) - scope-0  
|   |  
|   |  
|  
2013-03-19 18:19:40,210 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCCompiler - File concatenation threshold: 100 optimistic? false  
2013-03-19 18:19:40,255 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1  
2013-03-19 18:19:40,255 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1  
#-----  
# Map Reduce Plan  
#-----  
MapReduce node scope-6  
Map Plan  
A: Store(fakefile:org.apache.pig.builtin.PigStorage) - scope-5  
|  
|---A: New For Each(false)[bag] - scope-4  
|   |  
|   | Cast[chararray] - scope-2  
|   | |  
|   | |---Project[bytearray][0] - scope-1  
|   | |  
|   | |---A: Load(/home/kma041000/pig/input:org.apache.pig.builtin.PigStorage) - scope-0-----  
Global sort: false  
-----  
grunt>
```



# Debug: illustrate

---

- `grunt> illustrate A;`
  - -----
  - | A | line:chararray |
  - -----
  - | | word count in pig tutorial |
  - -----



# Pig WordCount – Batch mode

---

- Script file name: wordcount.pig
- Contains:

-----

```
A = load '/home/kma041000/pig/input';
```

```
B = foreach A generate flatten(TOKENIZE((chararray)$0)) as  
    word;
```

```
C = group B by word;
```

```
D = foreach C generate COUNT(B), group;
```

```
/* rm '/home/kma041000/pig/output'; */
```

```
store D into '/home/kma041000/pig/output';
```

-----



# Pig WordCount – Batch mode

- To run wordcount.pig
- Batch mode:
- {cs6360:~/BigData/Pig} pig -x mapreduce wordcount.pig
- .
- . *Hadoop map/reduce is running...*
- .
- {cs6360:~/BigData/Pig} hadoop fs -cat /home/kma041000/pig/output/part-r-00000
- 4 in
- 2 for
- 4 pig
- 4 2012
- 2 word
- 2 count
- 4 school
- 4 summer
- 2 indiana
- 4 tutorial



# Pig WordCount – Interactive mode

---

- {cs6360:~/BigData/Pig} pig
- grunt> A = load '/home/kma041000/pig/input';
- grunt> B = foreach A generate flatten(TOKENIZE((chararray) \$0)) as word;
- grunt> C = group B by word;
- grunt> D = foreach C generate COUNT(B), group;
- grunt> rm '/home/kma041000/pig/output';
- grunt> dump D; /\* to see output in terminal \*/
- grunt> store D into '/home/kma041000/pig/output'; /\* to part-r-00000 file \*/



# Pig WordCount – Interactive mode

---

- `grunt> cat /home/kma041000/pig/output/part-r-00000`
- 4 in
- 2 for
- 4 pig
- 4 2012
- 2 word
- 2 count
- 4 school
- 4 summer
- 2 indiana
- 4 tutorial
- Or
- `grunt> quit;`
- `{cs6360:~/BigData/Pig} hadoop fs -cat /home/kma041000/pig/output/part-r-00000`



# Pig UDF

---

- Pig provides extensive support for user-defined functions (UDFs) as a way to specify custom processing.
- Functions can be a part of almost every operator in Pig.



# Sample UDF Function

---

```
import java.io.IOException;
import java.util.logging.Level;
import java.util.logging.Logger;
import org.apache.pig.EvalFunc;
import org.apache.pig.backend.executionengine.ExecException;
import org.apache.pig.data.Tuple;
```

```
public class ConvertToUpper extends EvalFunc <String> {
```

```
    @Override
```

```
    public String exec(Tuple input) {
```

```
        try {
```

```
            if (input == null || input.size() == 0) {
```

```
                return null;
```

```
            }
```

```
            String str = (String) input.get(0);
```

```
            return str.toUpperCase();
```

```
        } catch (ExecException ex) {
```

```
            System.out.println("Error: " + ex.toString());
```

```
        }
```

```
        return null;
```

```
    }
```

```
}
```



# Compile & Run Jar

---

Compiling java file:

```
{cs6360:~} mkdir PIG_UDF
```

```
{cs6360:~} cd PIG_UDF
```

Copy the `ConvertToUpper.java` file to this directory.

```
{cs6360:~/PIG_UDF} javac -cp /usr/local/pig-0.10.1/pig-0.10.1.jar  
    ConvertToUpper.java
```

Create the jar file

```
{cs6360:~/PIG_UDF} jar -cf pig_udf.jar .
```



# Create Pig Script

---

Create a script `pig_script.pig` in `PIG_UDF` folder. The script has the following:

```
REGISTER /people/cs/m/mxs121731/HW_TA/PIG_UDF/pig_udf.jar;
```

```
A = LOAD '/HW_3_Data/movies_new' using PigStorage(';') as (MOVIEID:  
    chararray, TITLE: chararray, GENRE: chararray);
```

```
B = FOREACH A GENERATE MOVIEID, ConvertToUpper(TITLE), GENRE;
```

```
C = LIMIT B 10;
```

```
DUMP C;
```

You can also do it in interactive mode (using grunt shell)

NB: The location of the `pig_udf.jar` file should be changed according to your location



# Run Pig Script

---

Running pig\_script.pig

```
{cs6360:~/PIG_UDF} pig -x mapreduce pig_script.pig
```

## Output

```
(1,TOY STORY (1995),Animation|Children's|Comedy)
(2,JUMANJI (1995),Adventure|Children's|Fantasy)
(3,GRUMPIER OLD MEN (1995),Comedy|Romance)
(4,WAITING TO EXHALE (1995),Comedy|Drama)
(5,FATHER OF THE BRIDE PART II (1995),Comedy)
(6,HEAT (1995),Action|Crime|Thriller)
(7,SABRINA (1995),Comedy|Romance)
(8,TOM AND HUCK (1995),Adventure|Children's)
(9,SUDDEN DEATH (1995),Action)
(10,GOLDENEYE (1995),Action|Adventure|Thriller)
```





# Pig Example

---

Query: List the top 10 average rated movies in descending order.

```
A = load '/HW_3_Data/ratings_new ' using PigStorage(';') as (USERID:int,  
    MOVIEID:int, RATING:double, TIMESTAMP:chararray);  
B = group A by MOVIEID;  
C = foreach B generate group, AVG(A.RATING) as avgRating;  
D = order C by avgRating desc;  
E = limit D 10;  
dump E;
```

Output:

(3280,5.0)
(989,5.0)
(3656,5.0)
(1830,5.0)
(3607,5.0)
(787,5.0)
(3382,5.0)
(3172,5.0)
(3881,5.0)
(3233,5.0)



# References

---

<http://pig.apache.org/>

<http://pig.apache.org/docs/r0.10.0/basic.html#comparison>

<http://pig.apache.org/docs/r0.7.0/udf.html>

<http://www.bidn.com/blogs/cprice1979/ssas/4218/mmm-more-bacon-pig-user-defined-functions-udfs>



# Thank You

---