

HIVE

A warehouse solution
over
map-reduce framework

Dony Ang

Ashish Thusoo, Joydeep Sen Sarma, Namit Jain, Zheng Shao,
Prasad Chakka, Suresh Anthony, Hao Liu, Pete Wyckoff and Raghotham Murthy


overview

- background
- what is Hive
- Hive DB
- Hive architecture
- Hive datatypes
- hiveQL
- hive components
- execution flows
- compiler in details
- pros and cons
- conclusion

background

- Size of collected and analyzed datasets for business intelligence is growing rapidly, making traditional warehousing more \$\$\$
- Hadoop is a popular open source map-reduce as an alternative to store and process extremely large data sets on commodity hardware
- However, map reduce itself is very low-level and required developers to write custom code.

General Ecosystem of DW



what is hive ?

- Open-source DW solution built on top of Hadoop
- Support SQL-like declarative language called HiveQL which are compiled into map-reduce jobs executed on Hadoop
- Also support custom map-reduce script to be plugged into query.
- Includes a system catalog, Hive Metastore for query optimizations and data exploration

Hive Database

- Data Model
 - Tables
 - Analogous to tables in relational database
 - Each table has a corresponding HDFS dir
 - Data is serialized and stored in files within dir
 - Support external tables on data stored in HDFS, NFS or local directory.
 - Partitions
 - @table can have 1 or more partitions (1-level) which determine the distribution of data within subdirectories of table directory.

HIVE Database cont.

e.g : Table T under /wh/T and is partitioned on column
ds + ctry

For ds=20090101

ctry=US

Then data is stored within dir
/wh/T/ds=20090101/ctry=US

- Buckets
 - Data in each partition are divided into buckets based on hash of a column in the table. Each bucket is stored as a file in the partition directory.

HIVE datatype

- Support primitive column types
 - Integer
 - Floating point
 - Strings
 - Date
 - Boolean
- As well as nestable collections such as array or map
- User can also define their own type programmatically

Data Units

▶ Databases.

▶ Tables.

▶ Partitions.

▶ Buckets (or Clusters)..

Type System

▶ Primitive types

- Integers: TINYINT, SMALLINT, INT, BIGINT.
- Boolean: BOOLEAN.
- Floating point numbers: FLOAT, DOUBLE .
- String: STRING.

▶ Complex types

- Structs: {a INT; b INT}.
- Maps: M['group'].
- Arrays: ['a', 'b', 'c'], A[1] returns 'b'.

Examples – DDL Operations

▶ **CREATE TABLE** sample (foo INT, bar STRING) **PARTITIONED BY** (ds STRING);

▶ **SHOW TABLES** '.*s';

▶ **DESCRIBE** sample;

▶ **ALTER TABLE** sample **ADD COLUMNS** (new_col INT);

▶ **DROP TABLE** sample;

Examples – DML Operations

▶ **LOAD DATA LOCAL INPATH './sample.txt'**
OVERWRITE INTO TABLE sample
PARTITION (ds='2012-02-24');

▶ **LOAD DATA INPATH**
'/user/falvariz/hive/sample.txt' OVERWRITE
INTO TABLE sample PARTITION (ds='2012-
02-24');

SELECTS and FILTERS

▶ **SELECT** foo **FROM** sample **WHERE**
ds='2012-02-24';

▶ **INSERT OVERWRITE DIRECTORY**
'/tmp/hdfs_out' **SELECT** * **FROM** sample
WHERE ds='2012-02-24';

▶ **INSERT OVERWRITE LOCAL DIRECTORY**
'/tmp/hive-sample-out' **SELECT** * **FROM**
sample;

hiveQL

- Support SQL-like query language called HiveQL for select, join, aggregate, union all and sub-query in the from clause
- Support DDL stmt such as CREATE table with serialization format, partitioning and bucketing columns
- Command to load data from external sources and INSERT into HIVE tables.

```
LOAD DATA LOCAL INPATH '/logs/status_updates'  
INTO TABLE status_updates PARTITION (ds='2009-03-20')
```
- DO NOT support UPDATE and DELETE

hiveQL cont.

- Support multi-table INSERT

```
FROM (SELECT a.status, b.schoold, b.gender
      FROM status_updates a JOIN profiles b
      ON (a.userid = b.userid)
      and a.ds='2009-03-20')
      subq1
INSERT OVERWRITE TABLE gender_summary PARTITION (ds='2009-03-20')
  SELECT subq1.gender, COUNT(1) GROUP BY subq1.gender
INSERT OVERWRITE TABLE school_summary PARTITION (ds='009-03-20')
  SELECT subq.school, COUNT(1) GROUP BY subq1.school
```

- Also support User-defined column transformation (UDF) and aggregation (UDAF) function written in Java

Aggregations and Groups

▶ **SELECT MAX(foo) FROM sample;**

▶ **SELECT ds, COUNT(*), SUM(foo) FROM
sample GROUP BY ds;**

▶ **FROM sample s INSERT OVERWRITE
TABLE bar SELECT s.bar, count(*) WHERE
s.foo > 0 GROUP BY s.bar;**

Join

```
CREATE TABLE customer (id INT,name STRING,address STRING)  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '#';  
CREATE TABLE order_cust (id INT,cus_id INT,prod_id INT,price INT)  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

- ▶ **SELECT * FROM** customer c **JOIN** order_cust o **ON** (c.id=o.cus_id);
- ▶ **SELECT** c.id,c.name,c.address,ce.exp **FROM** customer c **JOIN** (**SELECT** cus_id,sum(price) AS exp **FROM** order_cust **GROUP BY** cus_id) ce **ON** (c.id=ce.cus_id);

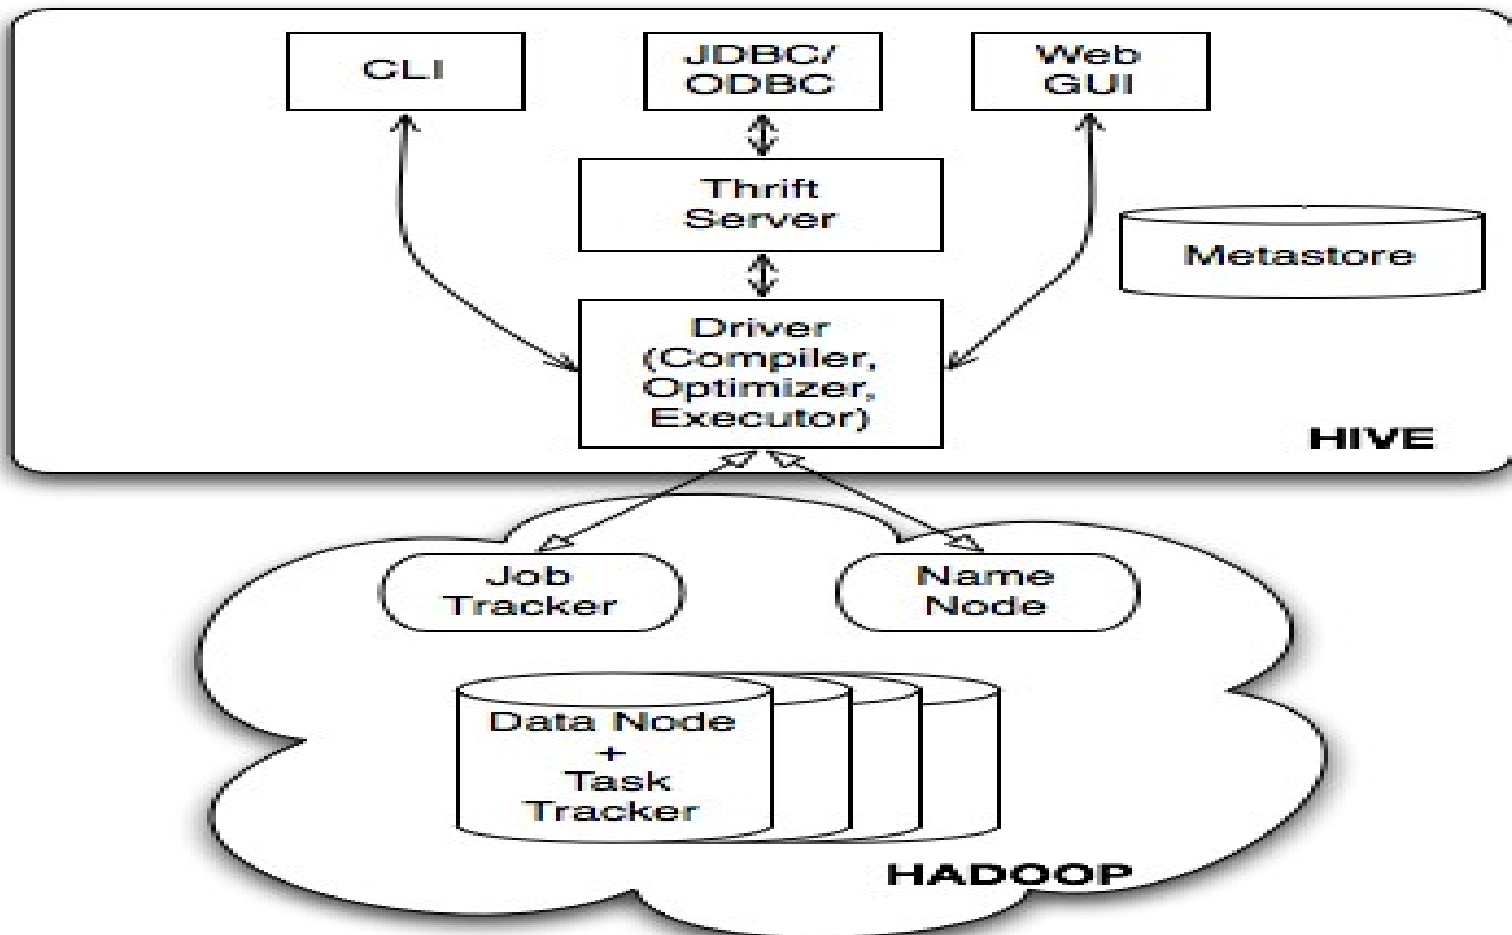
Multi table insert - Dynamic partition insert

```
FROM page_view_stg pvs
  INSERT OVERWRITE TABLE page_view PARTITION(dt='2008-06-08', country='US')
    SELECT pvs.viewTime, ... WHERE pvs.country = 'US'
  INSERT OVERWRITE TABLE page_view PARTITION(dt='2008-06-08', country='CA')
    SELECT pvs.viewTime, ... WHERE pvs.country = 'CA'
  INSERT OVERWRITE TABLE page_view PARTITION(dt='2008-06-08', country='UK')
    SELECT pvs.viewTime, ... WHERE pvs.country = 'UK';
```

```
FROM page_view_stg pvs
  INSERT OVERWRITE TABLE page_view PARTITION(dt='2008-06-08', country)
    SELECT pvs.viewTime, ...
```

<https://cwiki.apache.org/confluence/display/Hive/Tutorial#Tutorial-Dynamic-PartitionInsert>

HIVE Architecture



HIVE Components

- External Interfaces
 - User Interfaces both CLI and Web UI and API likes JDBC and ODBC.
- Hive Thrift Server
 - simple client API to execute HiveQL statements
- Metastore – system catalog
- Driver
 - Manages the lifecycle of HiveQL for compilation, optimization and execution.

Execution Flow