

# DISTRIBUTED INFORMATION SYSTEMS



The Pig Experience:  
Building High-Level Data  
flows on top of Map-Reduce


**VLDB paper**

**Source: Javeria Iqbal, Martin Theobald**

- Map-Reduce and the need for Pig Latin
- Pig Latin
- Compilation into Map-Reduce
- Optimization
- Future Work

# Data Processing Renaissance

---



- Internet companies swimming in data
  - TBs/day for Yahoo! Or Google!
  - PBs/day for FaceBook!
- Data analysis is “inner loop” of product innovation



# Data Warehousing ...?

---

## Scale

- High level declarative approach
- Little control over execution method

## Price

Prohibitively expensive at web scale

- Up to \$200K/TB

## SQL

Often not scalable enough

# Map-Reduce

---

- **Map** : Performs filtering
- **Reduce** : Performs the aggregation
- These are two high level declarative primitives to enable parallel processing
- **BUT** no complex Database Operations  
e.g. Joins



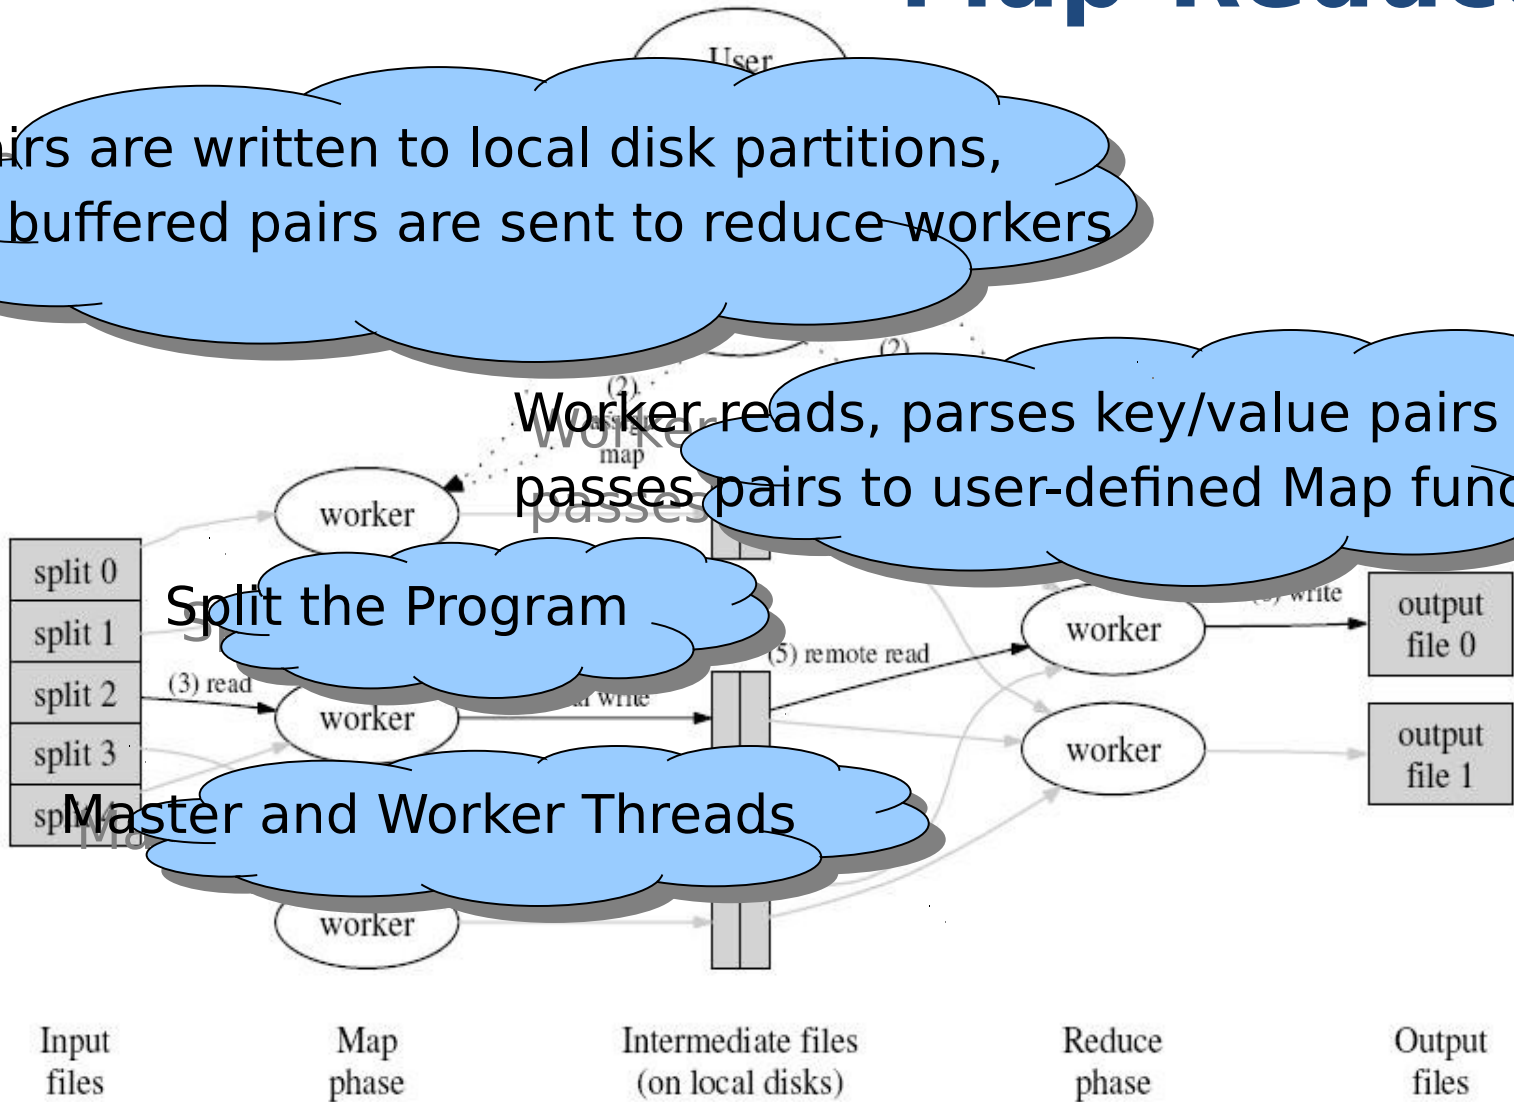
# Execution Overview of Map-Reduce

erred pairs are written to local disk partitions,  
tion of buffered pairs are sent to reduce workers

Worker reads, parses key/value pairs and  
passes pairs to user-defined Map function

Split the Program

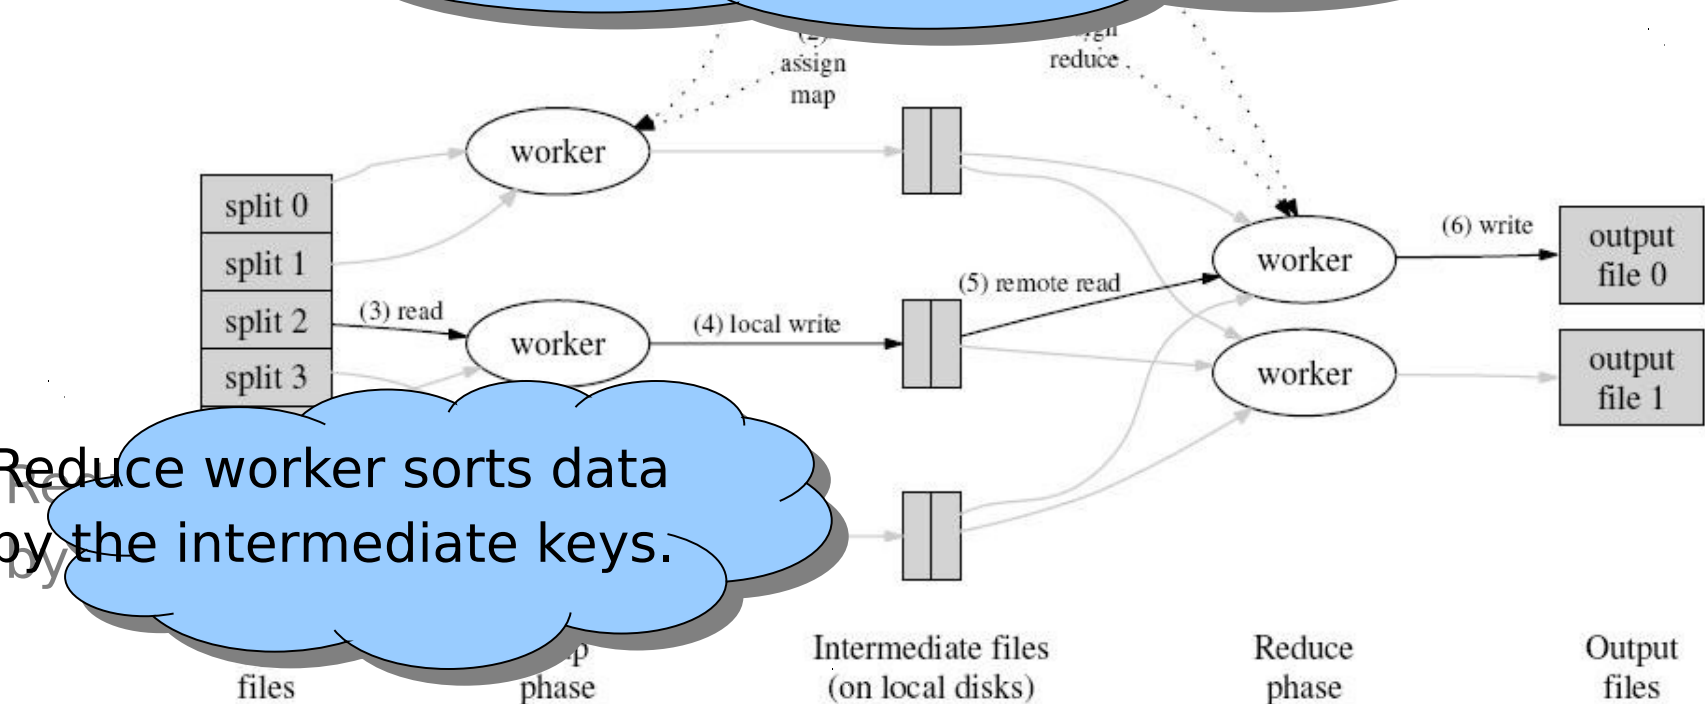
Master and Worker Threads





# Execution Overview of Map-Reduce

Unique keys, values are passed to user's Reduce function.  
Output is appended to the output file for this reduce partition.



Reduce worker sorts data  
by the intermediate keys.



# The Map-Reduce Appeal

---

Scale

- Scalable due to simpler design
- Explicit programming model

Price

- Only parallelizable operations
- Runs on cheap commodity hardware

Less Administration

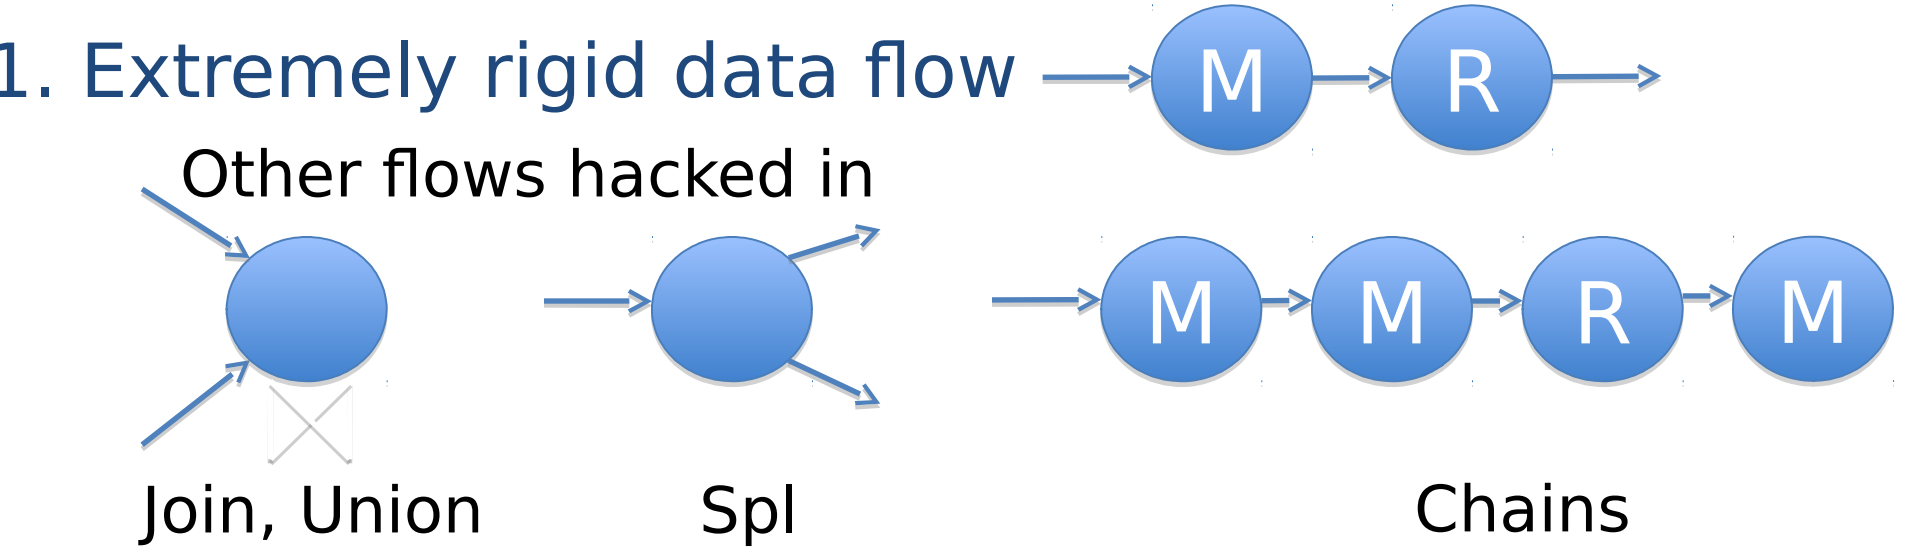
~~SQL~~

**Procedural Control**- a processing “pipe”





# Disadvantages



Common operations must be coded by hand  
Join, filter, projection, aggregates, sorting, distinct


Semantics hidden inside map-reduce functions  
Difficult to maintain, extend, and optimize

No combined processing of multiple Datasets  
Joins and other data processing operations

# Motivation

---

Need a high-level, general data flow language






# Enter Pig Latin

---

Need a high-level, general data flow language



- Map-Reduce and the need for Pig Latin
- Pig Latin
- Compilation into Map-Reduce
- Optimization
- Future Work



# Pig Latin: Data Types

---

- Rich and Simple Data Model

## Simple Types:

int, long, double, chararray, bytearray

## Complex Types:

- Atom: String or Number e.g. ('apple')
- Tuple: Collection of fields e.g. ('apple', 'mango')
- Bag: Collection of tuples

```
{ ('apple' , 'mango')  
  ('apple', ('red' , 'yellow'))  
}
```


- Map: Key, Value Pair



# Example: Data Model

---

- **Atom:** contains Single atomic value
- **Tuple:** sequence of fields
- **Bag:** collection of tuple with possible duplicates





# Pig Latin: Input/Output Data

---

## Input:

```
queries = LOAD `query_log.txt`  
USING myLoad()  
AS (userId, queryString, timestamp);
```

## Output:

```
STORE query_revenues INTO  
`myoutput`  
USING myStore();
```



# Pig Latin: General Syntax

---

- Discarding Unwanted Data:  
FILTER
- Comparison operators such  
as ==, eq, !=, neq
- Logical connectors AND, OR, NOT





# Pig Latin: Expression Table

$$t = \left( \text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

Let fields of tuple  $t$  be called  $f1$ ,  $f2$ ,  $f3$

Expression Type	Example	Value for $t$
Constant	'bob'	Independent of $t$
Field by position	$\$0$	'alice'
Field by name	$f3$	'age' $\rightarrow$ 20
Projection	$f2.\$0$	$\left\{ \begin{array}{l} (\text{'lakers'}) \\ (\text{'iPod'}) \end{array} \right\}$
Map Lookup	$f3\#\text{'age'}$	20
Function Evaluation	$\text{SUM}(f2.\$1)$	$1 + 2 = 3$
Conditional Expression	$f3\#\text{'age'} > 18?$ 'adult': 'minor'	'adult'
Flattening	$\text{FLATTEN}(f2)$	'lakers', 1 'iPod', 2

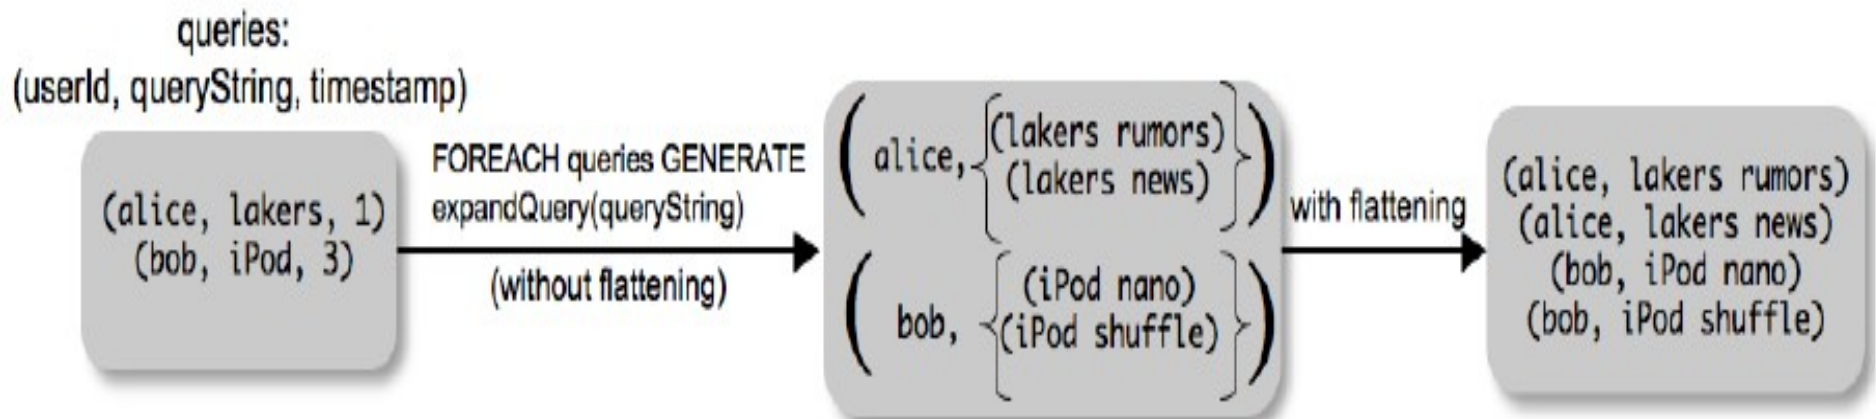
# Pig Latin: FOREACH with Flatten

---

```
expanded_queries = FOREACH
queries GENERATE userId,
expandQuery(queryString);
```

-----

```
expanded_queries = FOREACH
queries GENERATE userId,
```





# Pig Latin: COGROUP

---

- **Getting Related Data Together:**  
COGROUP

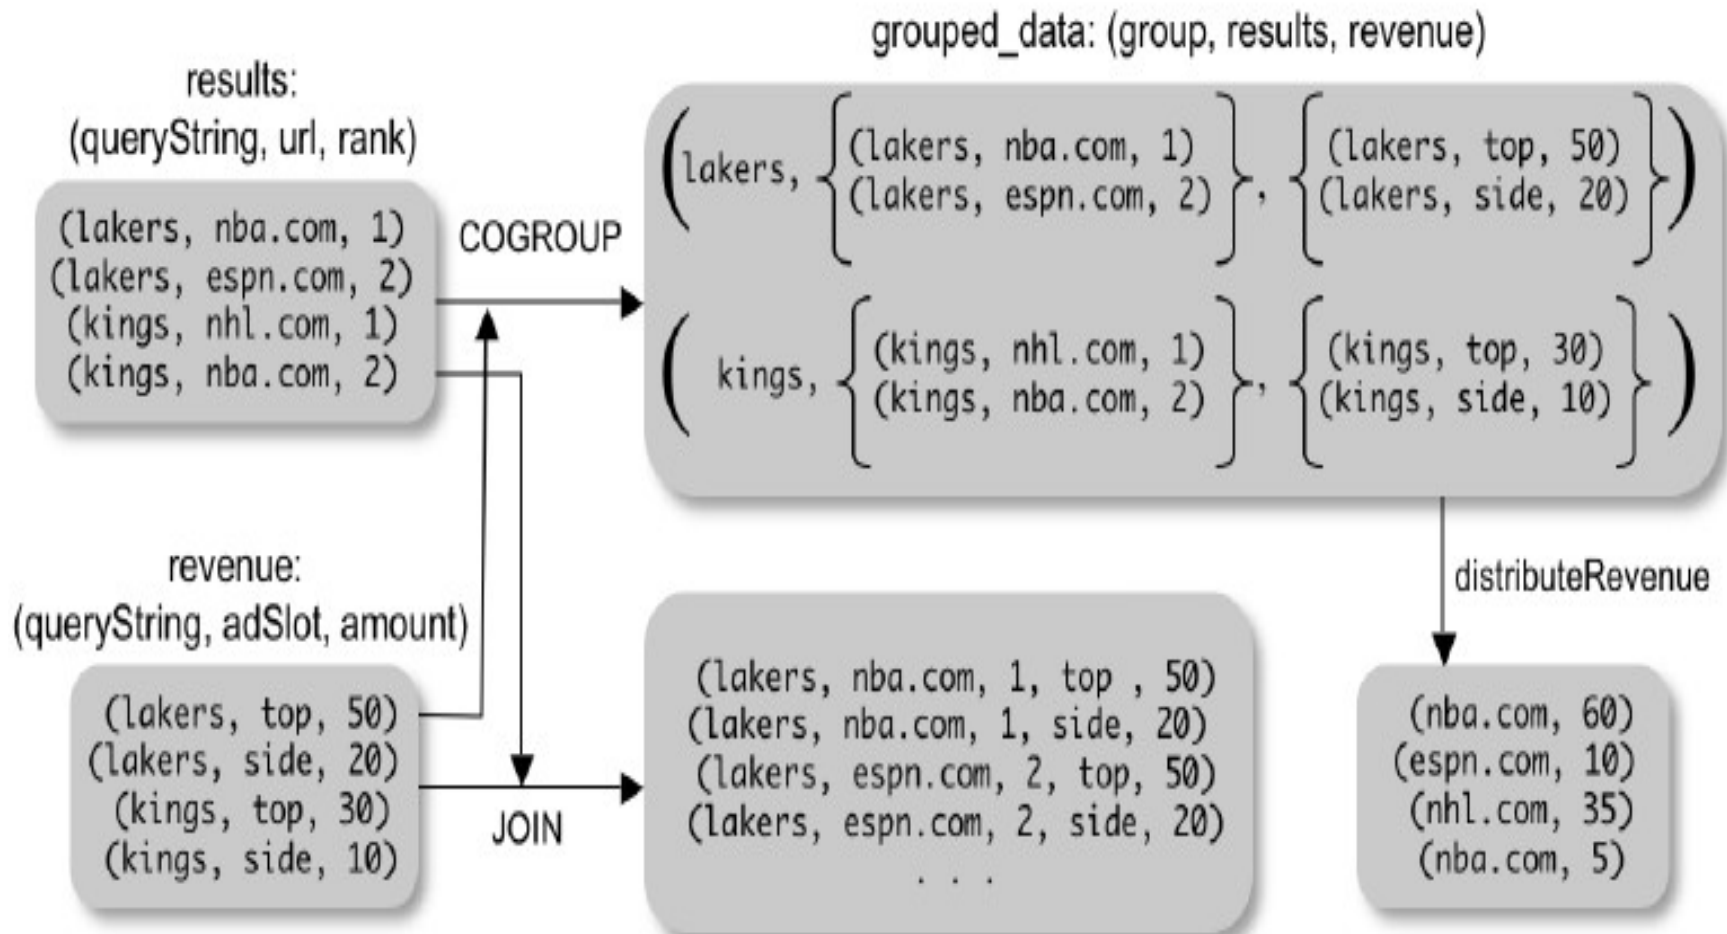
Suppose we have two data sets

result: (queryString, url, position)

revenue: (queryString, adSlot, amount)

```
grouped_data = COGROUP result BY  
queryString,                revenue BY  
queryString;
```

# Pig Latin: COGROUP vs. JOIN





# Pig Latin: Map-Reduce

---

- Map-Reduce in Pig Latin

```
map_result = FOREACH input GENERATE  
              FLATTEN(map(*));
```

```
key_group = GROUP map_result BY $0;
```

```
output = FOREACH key_group GENERATE  
              reduce(*);
```



# Pig Latin: Other Commands

---

- UNION : Returns the union of two or more bags
- CROSS: Returns the cross product
- ORDER: Orders a bag by the specified field(s)
- DISTINCT: Eliminates duplicate tuple in a bag



# Pig Latin: Nested Operations

---

```
grouped_revenue = GROUP revenue BY
  queryString;
query_revenues = FOREACH
  grouped_revenue {
top_slot = FILTER revenue BY
  adSlot eq `top';
GENERATE queryString,
SUM(top_slot.amount),
SUM(revenue.amount);
};
```



# Pig Pen: Screen Shot

Operators

LOADGROUPCOGROUPFILTERFOREACHORDER

= LOAD  USING 

Default

 AS (  )

[Generate Query](#)

<pre>visits = LOAD 'visits.txt' AS (user, url, time);  pages = LOAD 'pages.txt' AS (url, pagerank);  v_p = JOIN visits BY url, pages BY url;  users = GROUP v_p BY user;  useravg = FOREACH users GENERATE group, AVG(v_p.pagerank) AS avgpr;  answer = FILTER useravg BY avgpr &gt; '0.5';</pre>	<pre>visits:  (Amy, cnn.com, 8am)         (Amy, frogs.com, 9am)         (Fred, snails.com, 11am)  pages:  (cnn.com, 0.8)         (frogs.com, 0.8)         (snails.com, 0.3)  v_p:    (Amy, cnn.com, 8am, cnn.com, 0.8)         (Amy, frogs.com, 9am, frogs.com, 0.8)         (Fred, snails.com, 11am, snails.com, 0.3)  users:  (Amy, { (Amy, cnn.com, 8am, cnn.com, 0.8),                (Amy, frogs.com, 9am, frogs.com, 0.8) })         (Fred, { (Fred, snails.com, 11am, snails.com, 0.3) })  useravg: (Amy, 0.8)          (Fred, 0.3)  answer: (Amy, 0.8)</pre>
---	--





# Pig Latin: Example 1

---

Suppose we have a table

urls: (url, category, pagerank)

Simple SQL query that finds,

For each sufficiently large category,  
the average pagerank of high-  
pagerank urls in that category

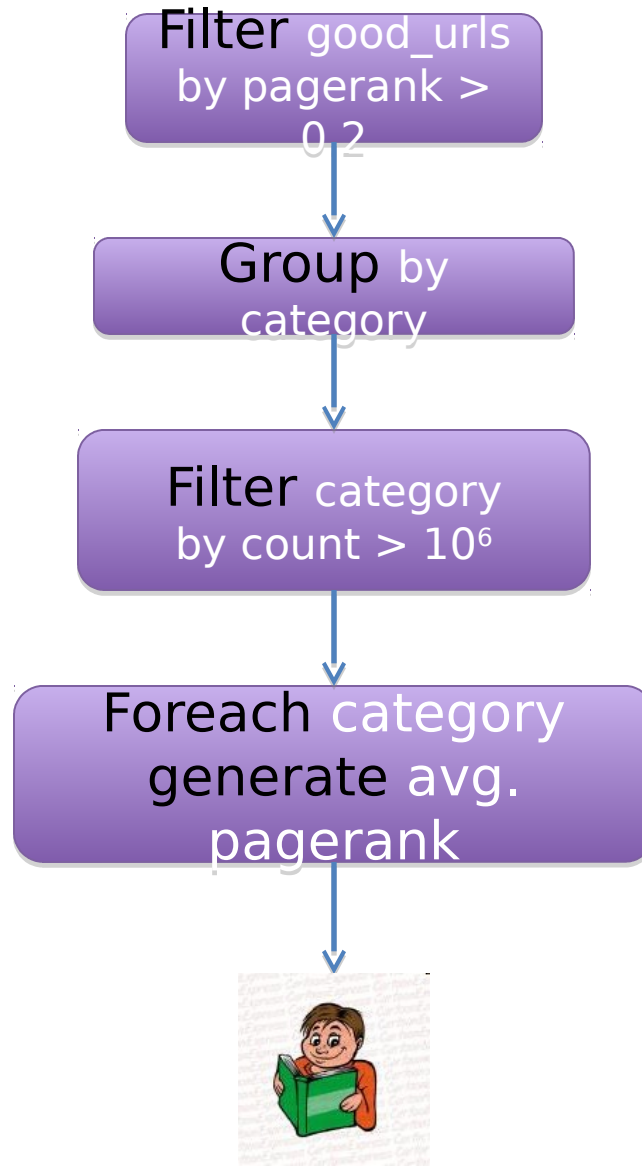
```
SELECT category, Avg(pagerank)
```

```
FROM urls WHERE pagerank > 0.2
```

```
GROUP BY category HAVING COUNT(*) > 106
```



# Data Flow





# Equivalent Pig Latin

---

- `good_urls` = FILTER urls BY pagerank > 0.2;
- `groups` = GROUP `good_urls` BY category;
- `big_groups` = FILTER `groups` BY  
COUNT(`good_urls`) >  $10^6$  ;
- `output` = FOREACH `big_groups` GENERATE  
category, AVG(`good_urls`.pagerank);



# Example 2: Data Analysis Task

Find the top 10 most visited pages in each category

Visits

User	Url	Time
Amy	cnn.com	8:00
Amy	bbc.com	10:00
Amy	flickr.com	10:05
Fred	cnn.com	12:00

⋮

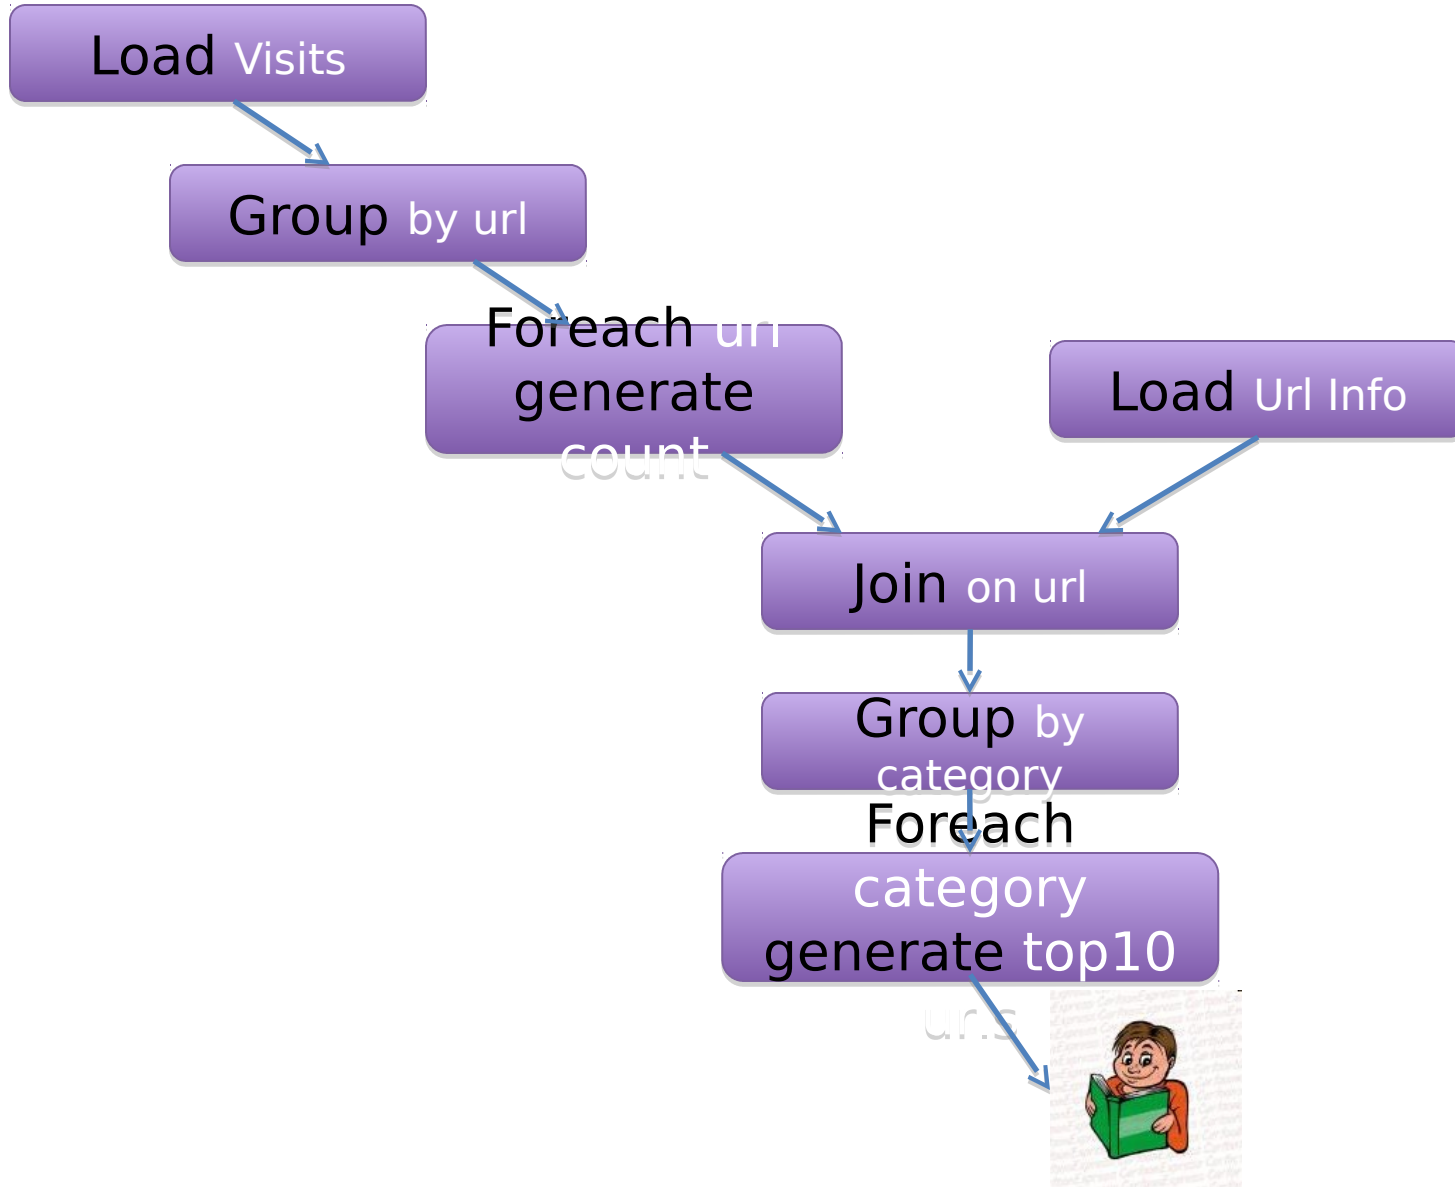
Url Info

Url	Category	PageRank
cnn.com	News	0.9
bbc.com	News	0.8
flickr.com	Photos	0.7
espn.com	Sports	0.9

⋮



# Data Flow





# Equivalent Pig Latin

---

```
visits          = load '/data/visits' as (user, url,
    time);
gVisits         = group visits by url;
visitCounts    = foreach gVisits generate url,
    count(visits);

urlInfo        = load '/data/urlInfo' as (url,
    category, pRank);
visitCounts    = join visitCounts by url, urlInfo by
    url;

gCategories    = group visitCounts by category;
topUrls = foreach gCategories generate
```



# Quick Start and Interoperability

---

visits = load '/data/visits' as (user, url, time);

gVisits = group visits by url;

visitCounts = foreach gVisits generate url, count(urlVisits);

urlInfo = load '/data/urlInfo' as (url, category);

Operates directly over files

visitCounts, urlInfo by url;

gCategories = group visitCounts by category;

topUrls = foreach gCategories generate top(visitCounts, 10);



# Quick Start and Interoperability

---

```
visits = load '/data/visits' as (user, url,  
time);
```

```
gVisits = group visits by url;
```

```
visitCounts = foreach gVisits generate url,  
count(urlVisits);
```

```
urlInfo = load '/data/urlInfo' as (url,  
category
```

Schemas optional;  
Can be assigned  
dynamically

```
visitCounts = join visitCounts by url, urlInfo by  
url;
```

```
gCategories = group visitCounts by category;
```

```
topUrls = foreach gCategories generate
```





# User-Code as a First-Class Citizen

visits  
time  
gVisits  
visitC  
count

User-defined functions (UDFs) can be used in every construct

- Load, Store
- Group, Filter, Foreach

s (user, url,  
erate url,

urlInfo = load '/data/urlInfo' as (url,  
category, pRank);

visitCounts = join visitCounts by url, urlInfo by  
url;

gCategories = group visitCounts by category;


topUrls = foreach gCategories generate



# Nested Data Model

---

- Pig Latin has a **fully nested data model** with:
  - Atomic values, tuples, bags (lists), and maps



- Avoids expensive joins



# Nested Data Model

Decouples grouping as an independent

operation

User	Url	Time
Amy	cnn.com	8:00
Amy	bbc.com	10:00
Amy	bbc.com	10:05
Fred	cnn.com	12:00

group by url



group	Visits		
cnn.com	Amy	cnn.com	8:00
	Fred	cnn.com	12:00
	Amy	bbc.com	10:00
			10:05

- Compact representation of data sets
- Powerful sequential operations
- Efficient Implementation (see paper)

I frankly like pig much better than SQL in some respects (group + optional flatten), I love nested data structures."

Ted Dunning

Chief Scientist, Veoh

## results

quer y	url	ran k
Laker s	nba.com	1
Laker s	espn.com	2
Kings	nhl.com	1
Kings	nba.com	2

## revenue

quer y	adSlot	amou nt
Laker s	top	50
Laker s	side	20
Kings	top	30
Kings	side	10

group	results			revenue		
Lakers	Laker s	nba.co m	1	Lakers	top	50
	Laker s	espn.co m	2	Lakers	side	20
Kings	Kings	nhl.co m	1	Kings	top	30
	Kings	nba.co m	2	Kings	side	10

Cross-product of the 2 bags would give natural join

# Pig Features

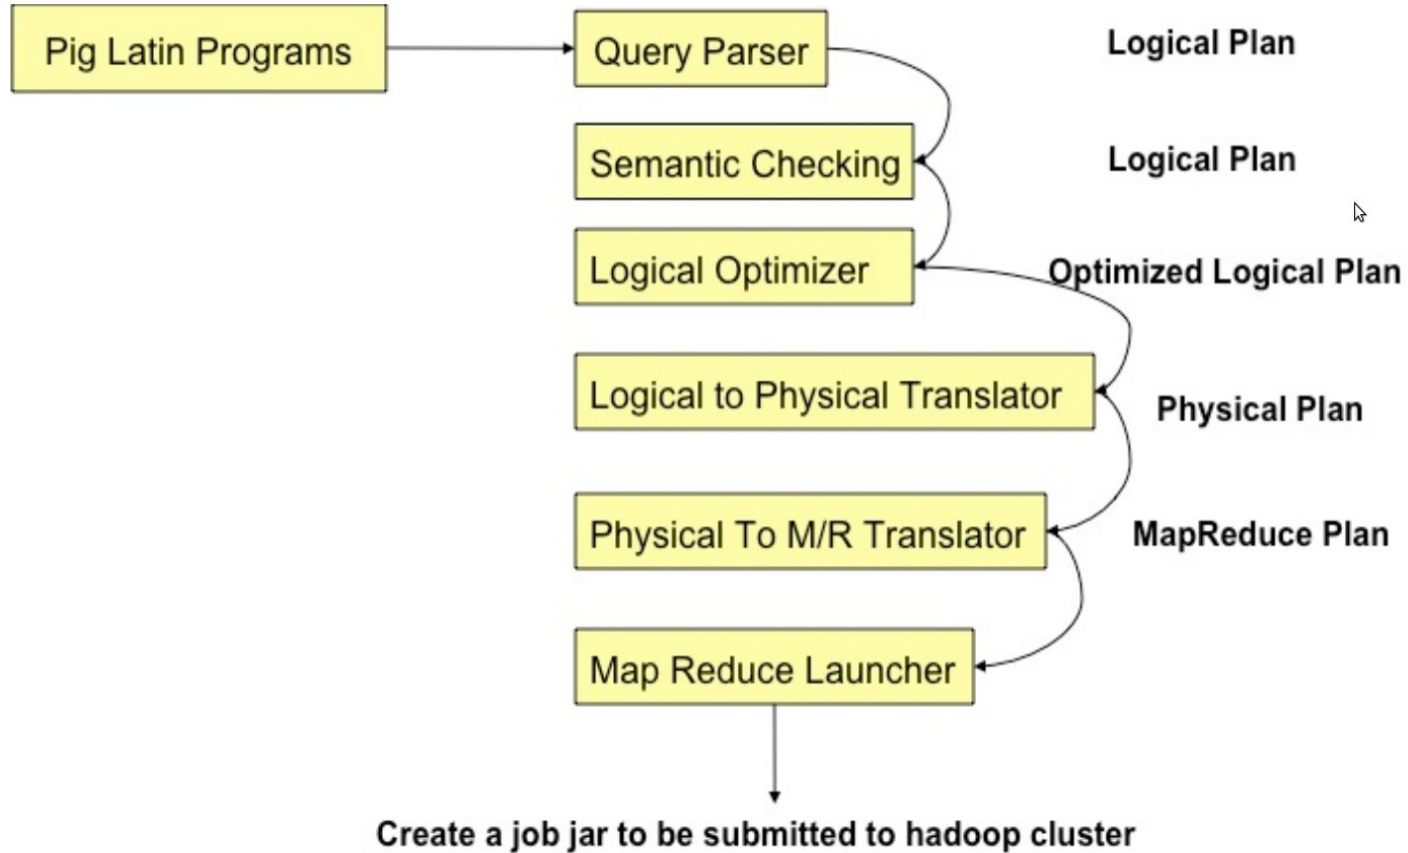
---

- Explicit Data Flow Language unlike SQL
- Low Level Procedural Language unlike Map-Reduce
- Quick Start & Interoperability
- Mode (Interactive Mode, Batch, Embedded)
- User Defined Functions
- Nested Data Model

- Map-Reduce and the need for Pig Latin
- Pig Latin
- **Compilation into Map-Reduce**
- Optimization
- Future Work



# Compilation





# Parsing

---


- Type checking with schema
- References verifying
- Logic plan generating
  - One-to-one fashion
  - Independent of execution platform
  - Limited optimization



# Y! Logical Plan

- Consists of DAG of Logical Operators as nodes and Data Flow represented as edges
  - Logical Operators contain list of i/p's o/p's and schema
- Logical operators
  - Aid in post parse stage checking (type checking)
  - Optimization
  - Translation to Physical Plan

```
a = load 'myfile';  
b = filter a by $0 > 5;  
store b into 'myfilteredfile';
```



# Physical Plan


---

- Layer to map Logical Plan to multiple back-ends, one such being M/R (Map Reduce)
  - Chance for code re-use if multiple back-ends share same operator
- Consists of operators which pig will run on the backend
- Currently most of the physical plan is placed as operators in the map reduce plan
- Logical to Physical Translation
  - 1:1 correspondence for most Logical operators
  - Except Cross, Distinct, Group, Co-group and Order



# Logic Plan

```
A=LOAD 'file1' AS (x, y, z);  
B=LOAD 'file2' AS (t, u, v);  
C=FILTER A by y > 0;  
D=JOIN C BY x, B BY u;  
E=GROUP D BY z;  
F=FOREACH E GENERATE  
  group, COUNT(D);  
STORE F INTO 'output';
```





# Physical Plan

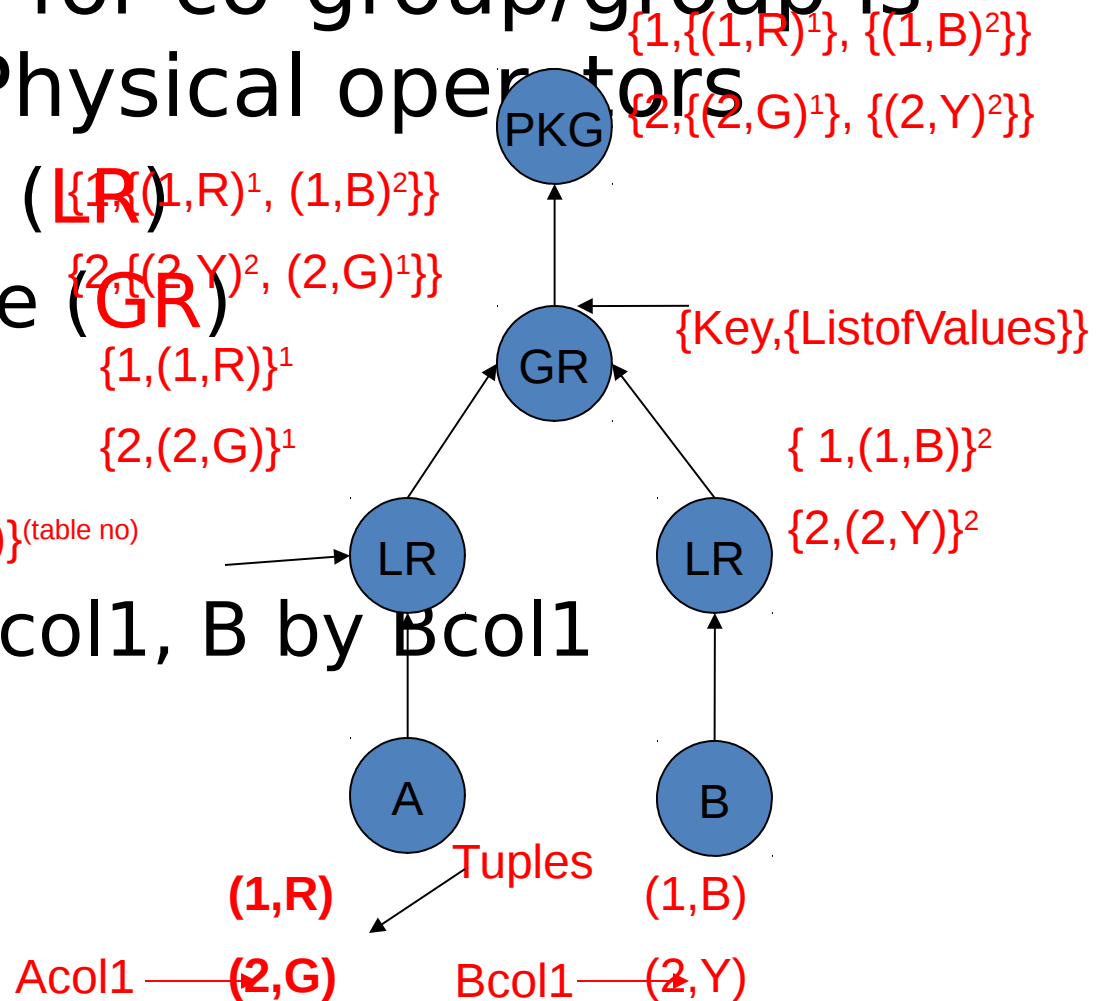
---

- 1:1 correspondence with most logical operators
- Except for:
  - DISTINCT
  - (CO)GROUP
  - JOIN
  - ORDER




# Logical to Physical Plan for Group operator

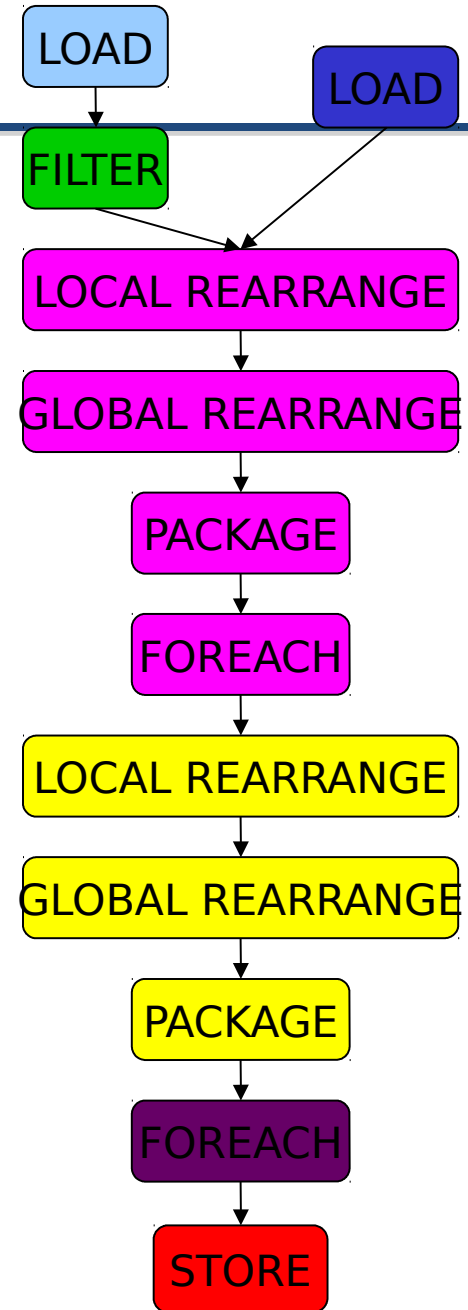

- Logical operator for co-group/group is converted to 3 Physical operators
  - Local Rearrange (**LR**)
  - Global Rearrange (**GR**)
  - Package (**PKG**)
- Example:
  - Co-group A by Acol1, B by Bcol1



# Y!Map Reduce Plan

- Physical to Map Reduce (M/R) Plan conversion happens through the MRCompiler
  - Converts a physical plan into a DAG of M/R operators
- Boundaries for M/R include cogroup/group, distinct, cross, order by, limit (in some cases)
  - Push all subsequent operators between cogroup to next cogroup into reduce
  - order by is implemented as 2 M/R jobs
- JobControlCompiler then uses the M/R plan to construct a JobControl object












# MapReduce Plan

---

- Determine MapReduce boundaries
  - GLOBAL REARRANGE
- Some operations are done by MapReduce framework
- Coalesce other operators into Map & Reduce stages
- Generate job jar file







# Pig Latin to Physical

A = LOAD 'file1' AS  
(x,y,z);

B = LOAD 'file2' AS  
(t,u,v);

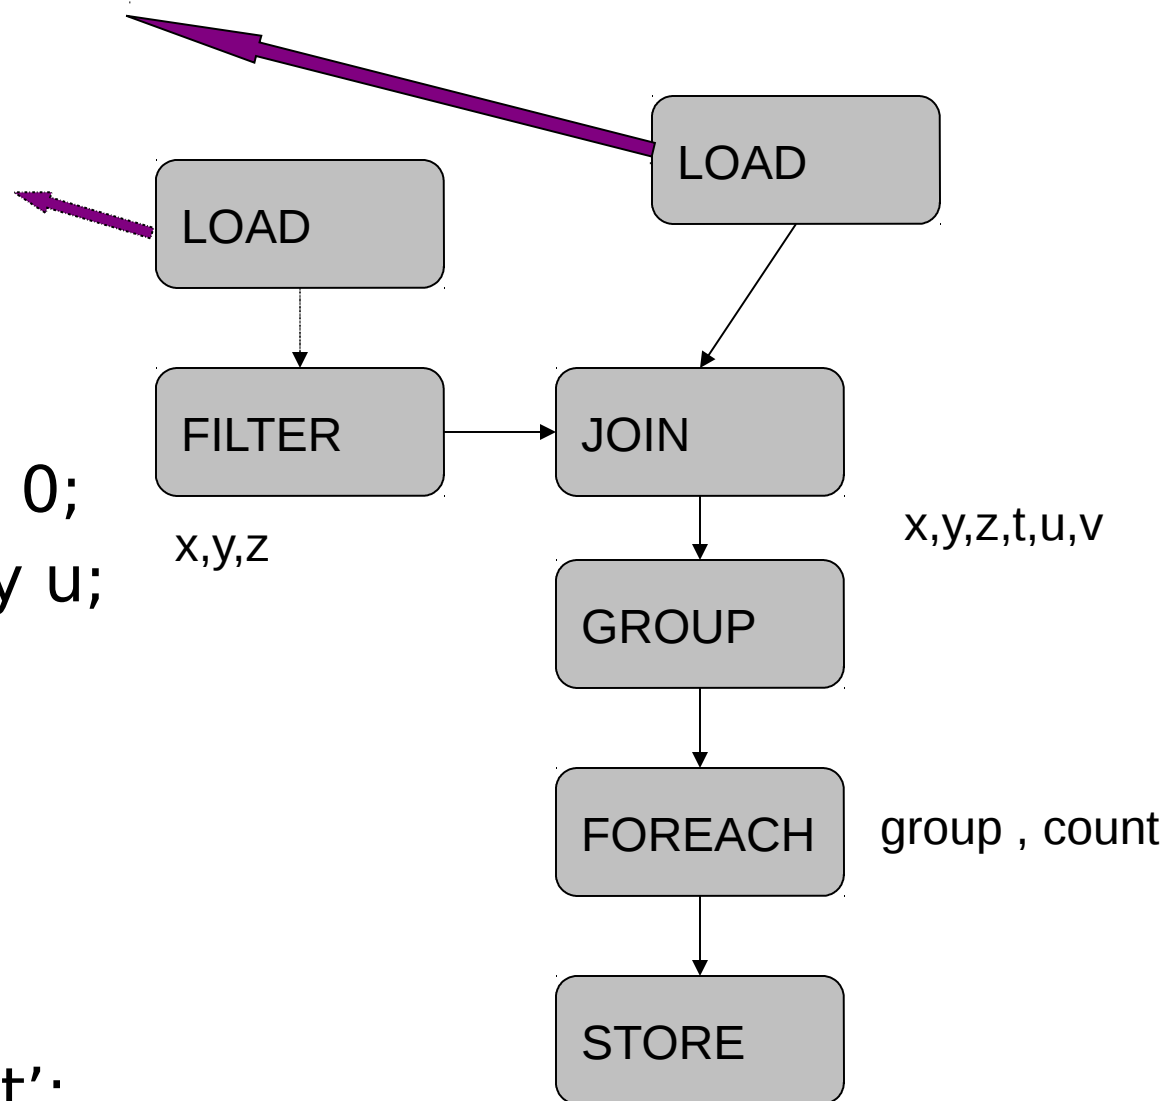
C = FILTER A by y > 0;

D = JOIN C by x, B by u;

E = GROUP D by z;

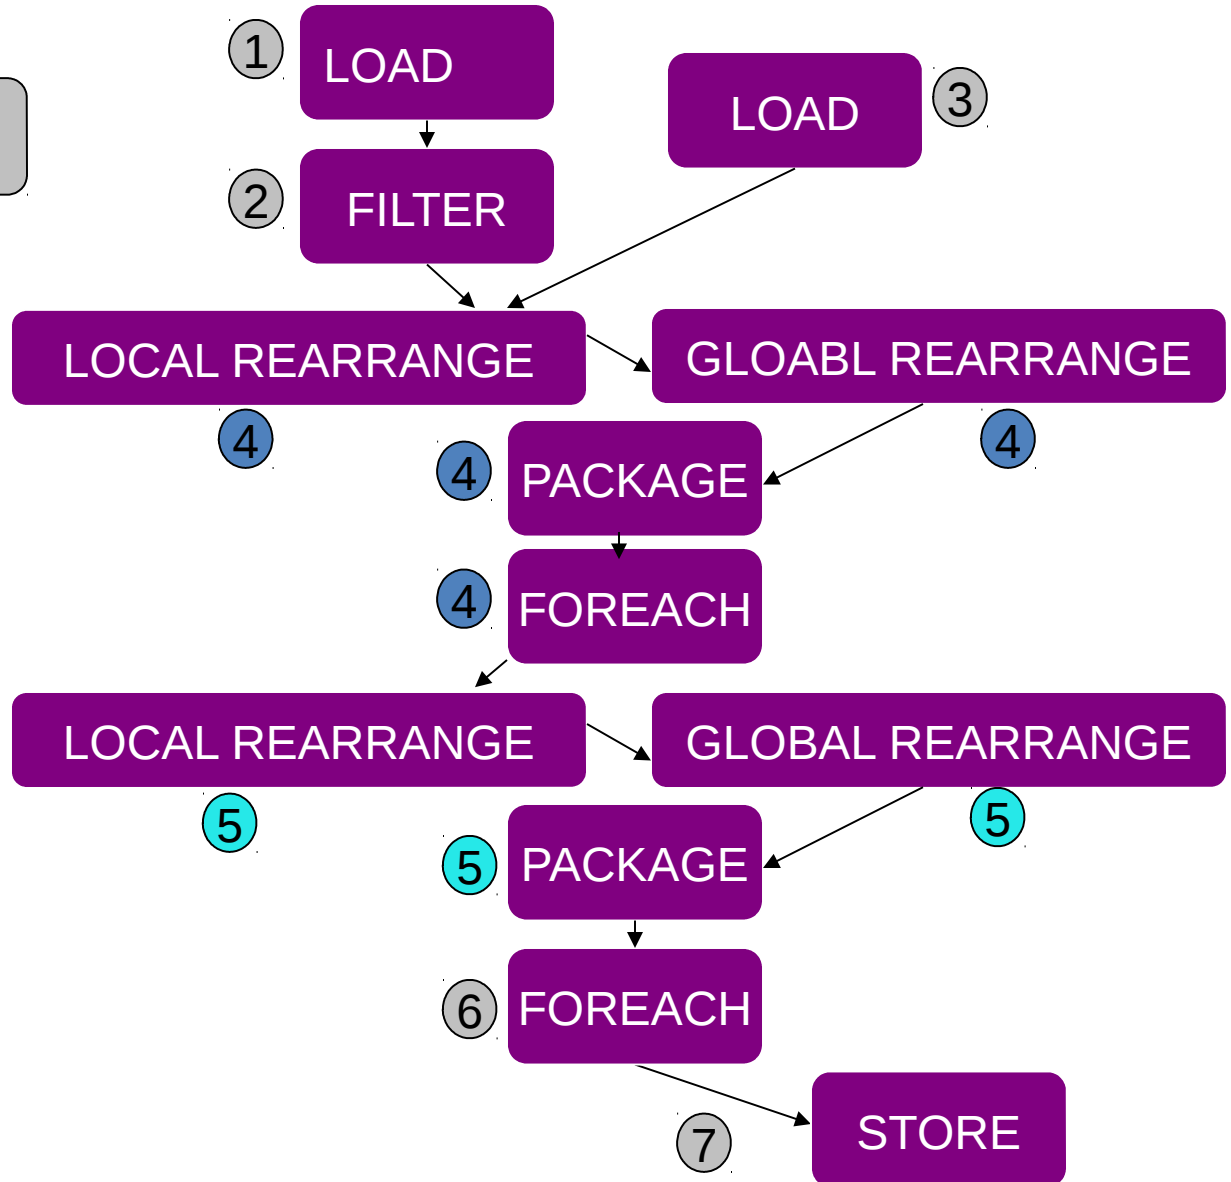
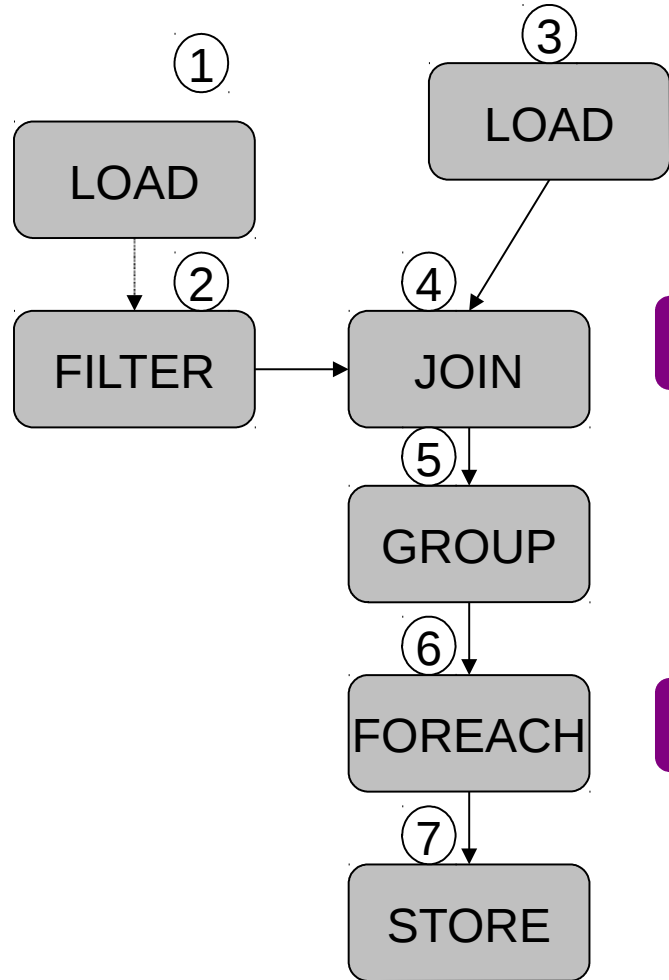
F = FOREACH E  
generate group,  
COUNT(D);

STORE F into 'output';







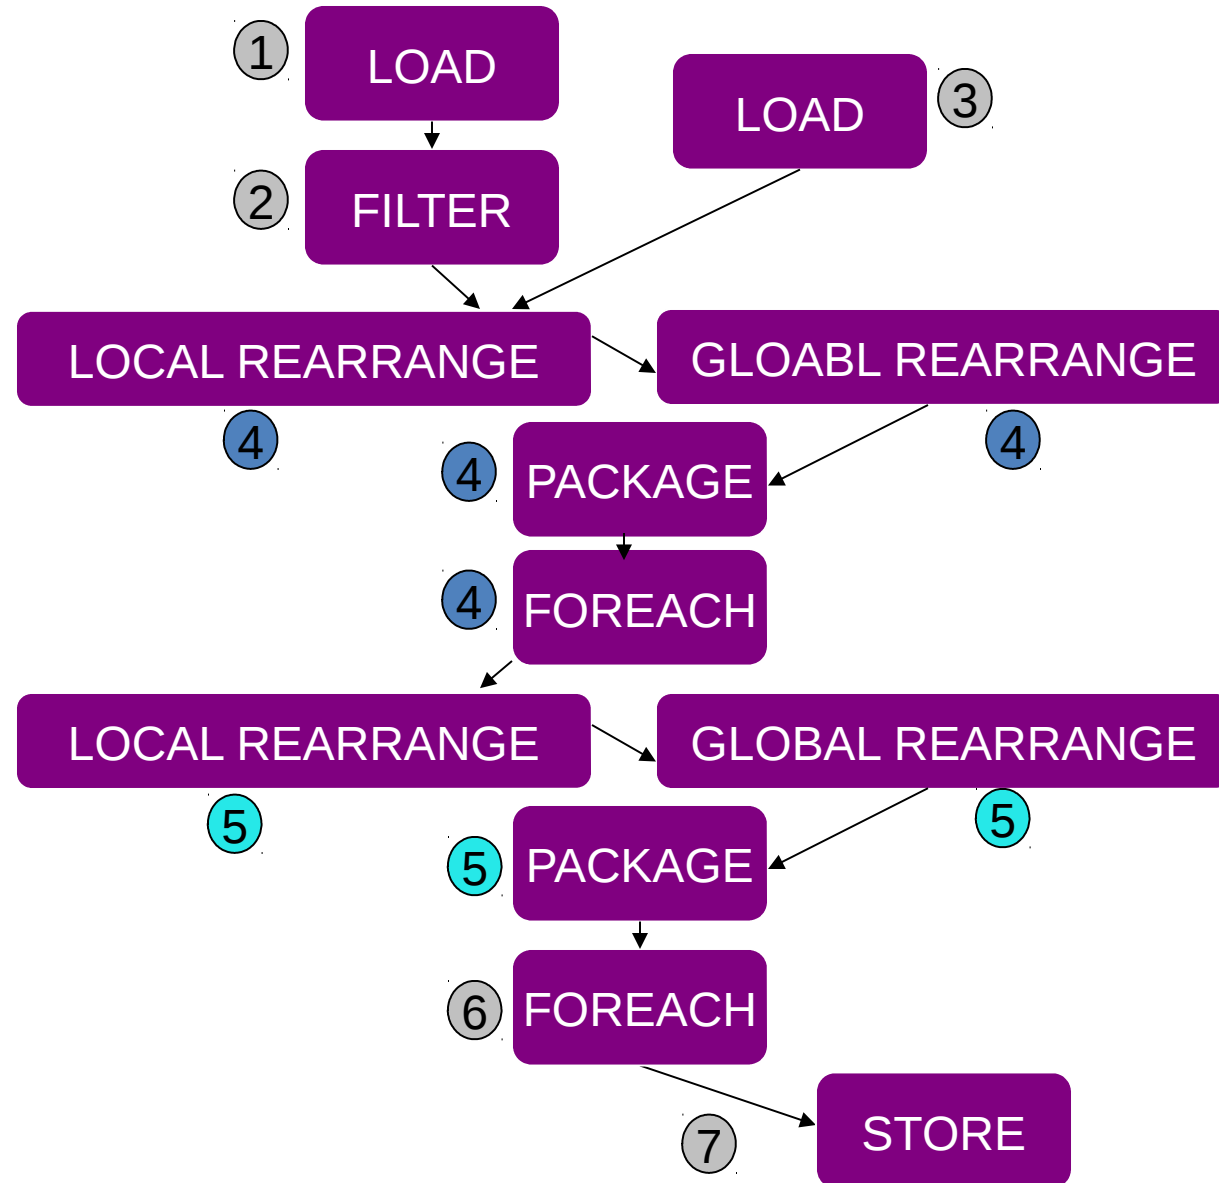


# Logical Plan to Physical Plan



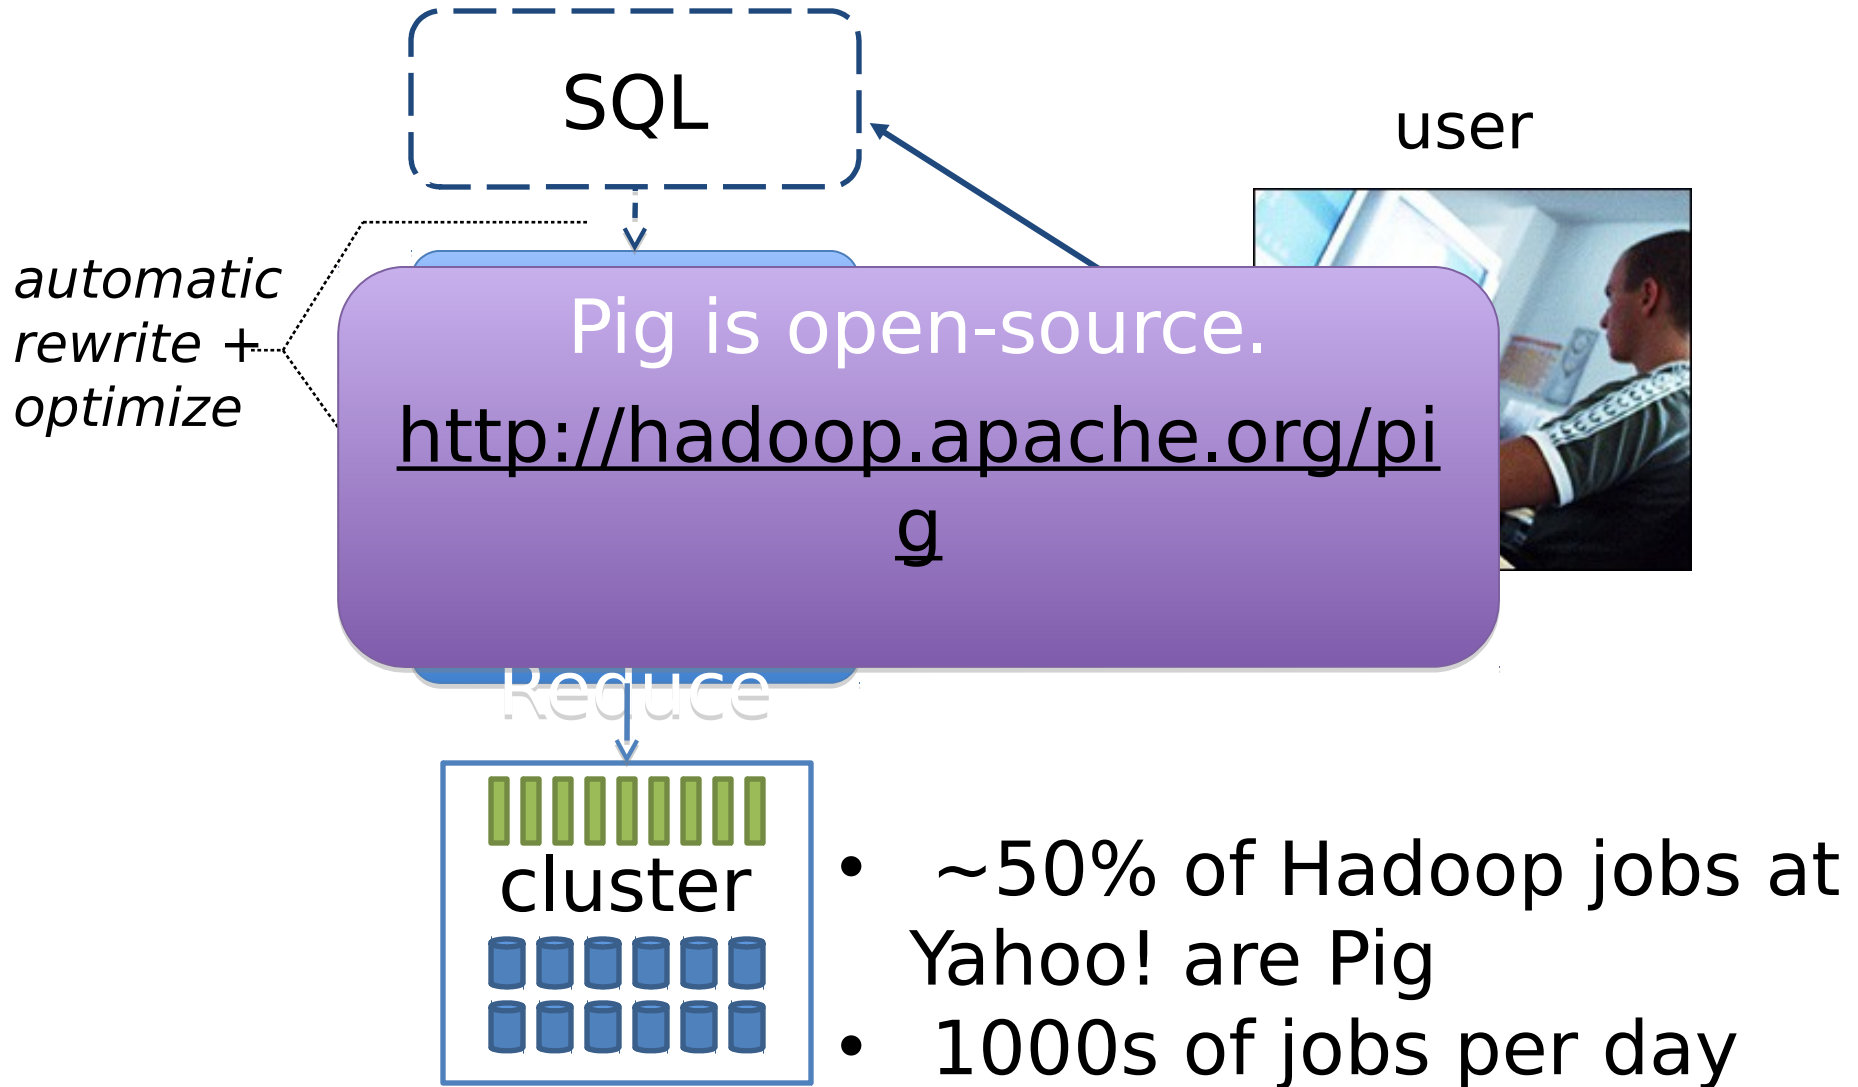


# Physical Plan to Map-Reduce Plan



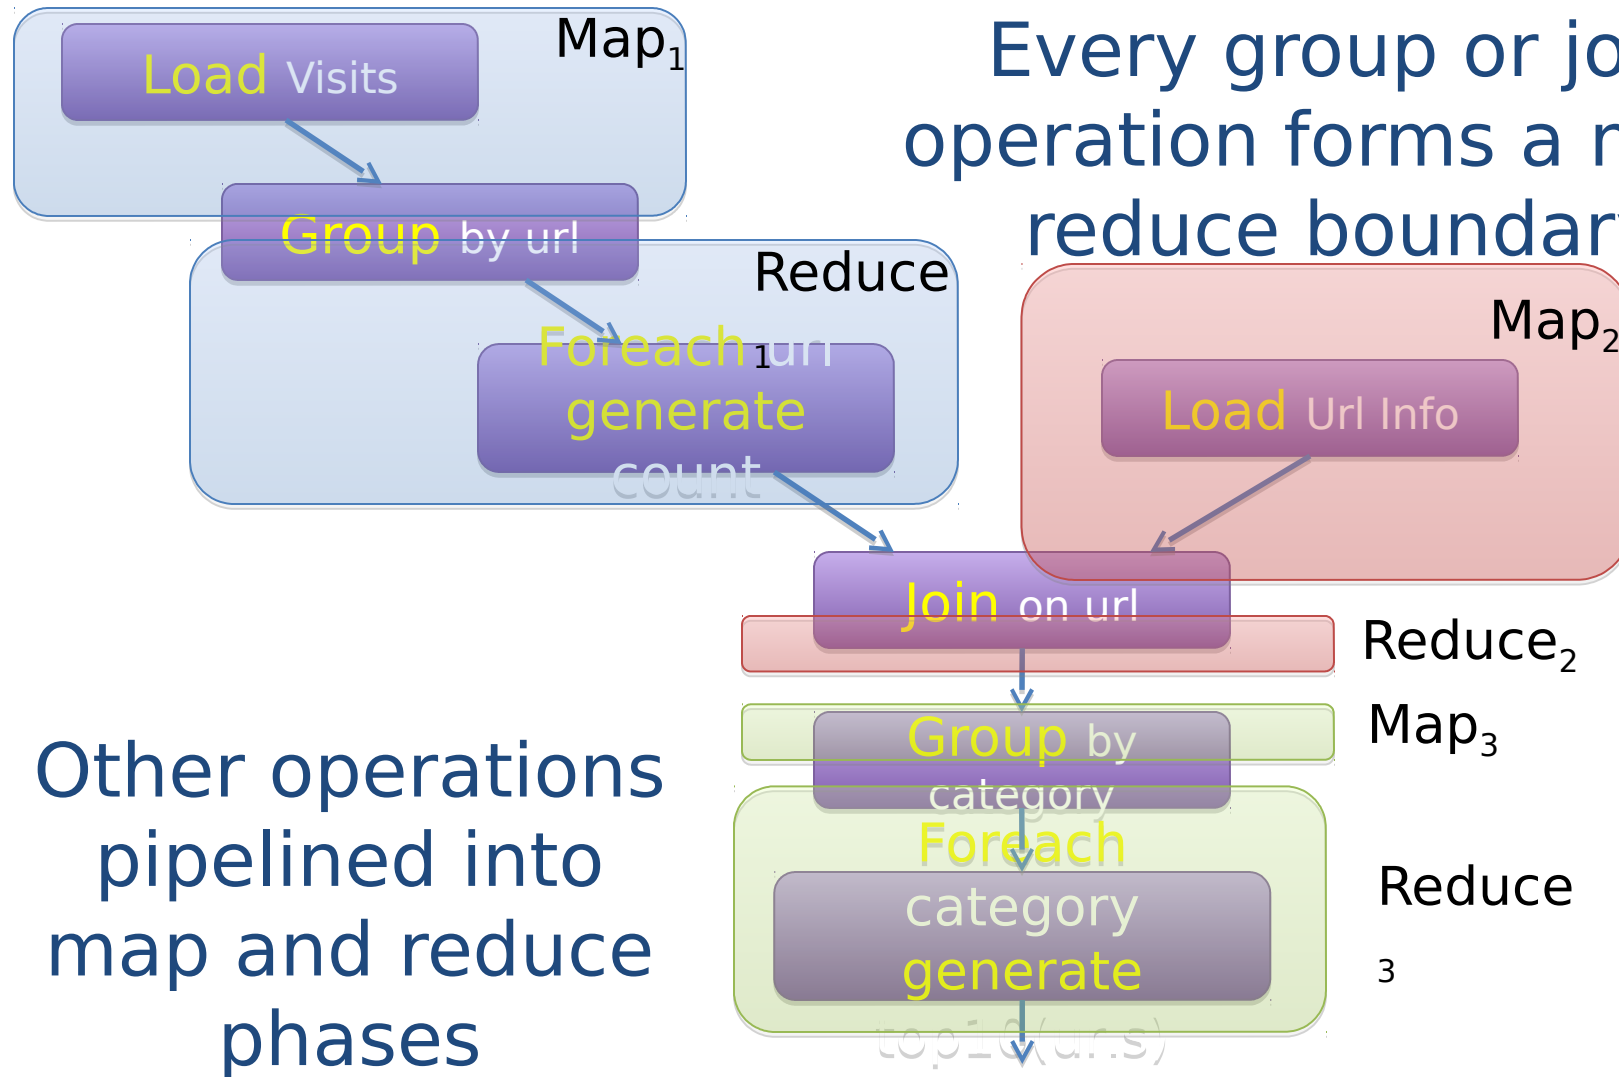


# Implementation






# Compilation into Map-Reduce







# Performance

## Pig Performance vs Map-Reduce





# Strong & Weak Points



The step-by-step method of creating a program in Pig is much cleaner and simpler to use than the single block method of SQL. It is easier to keep track of what your variables are, and where you are in the process of analyzing your data.

- **Scalability**

With the various interleaved clauses in SQL  
It is difficult to know what is actually happening sequentially.

- **Memory**

Jasmine Novak  
Engineer, Yahoo!

- **Processing**

- **Open Source**

- **Non Java Users**

David Cierniewicz  
Search Excellence, Yahoo!

- **Limited Optimization**



- Big demand for parallel data processing
  - Emerging tools that do not look like SQL DBMS
  - Programmers like dataflow pipes over static files
- Hence the excitement about Map-Pig Latin
  - Sweet spot between map-reduce and SQL
- But, Map-Reduce is too low-level and rigid